

## INTRODUCTION

This dataset describes web traffic on an anonymous webpage. It describes the origin of the traffic, the number of pages viewed, and more. It also contains the conversion rate, the target variable for converting traffic into desired actions (like purchases).

The features are as follows:

- Page Views: The number of pages the user viewed during the session
- Session Duration: The length of the session in minutes
- Bounce Rate: The percentage of visitors who left after visiting a single page. Exact calculation unknown.
- Traffic Source: The way the traffic originated, be it through organic search or through paid ads.
- Time on Page: The amount of time, in seconds, user spent on the specific page as the data was captured.
- Previous Visits: The number of times the current user has visited the page in the past.
- Conversion Rate: The percentage of users during the session that completed a desired interaction. Exact calculation unknown. Target variable.

The dataset was found on Kaggle here:

<https://www.kaggle.com/datasets/anthonytherrien/website-traffic>

The dataset contains 2000 rows and 7 features.

## LOADING DATA

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv('/content/website_wata.csv')

df.head()
```



	Page Views	Session Duration	Bounce Rate	Traffic Source	Time on Page	Previous Visits	Conversion Rate
0	5	11.051381	0.230652	Organic	3.890460	3	1.0
1	4	3.429316	0.391001	Social	8.478174	0	1.0
2	4	1.621052	0.397986	Organic	9.636170	2	1.0
3	5	3.629279	0.180458	Organic	2.071925	3	1.0
4	5	4.235843	0.291541	Paid	1.960654	5	1.0

## Exploratory Data Analysis

```
df.shape
```



```
(2000, 7)
```

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Page Views            2000 non-null   int64
1   Session Duration      2000 non-null   float64
2   Bounce Rate           2000 non-null   float64
3   Traffic Source        2000 non-null   object
4   Time on Page          2000 non-null   float64
5   Previous Visits       2000 non-null   int64
6   Conversion Rate       2000 non-null   float64
dtypes: float64(4), int64(2), object(1)
memory usage: 109.5+ KB
```

```
df.describe()
```



	Page Views	Session Duration	Bounce Rate	Time on Page	Previous Visits	Conversion Rate
<b>count</b>	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
<b>mean</b>	4.950500	3.022045	0.284767	4.027439	1.978500	0.982065
<b>std</b>	2.183903	3.104518	0.159781	2.887422	1.432852	0.065680
<b>min</b>	0.000000	0.003613	0.007868	0.068515	0.000000	0.343665
<b>25%</b>	3.000000	0.815828	0.161986	1.935037	1.000000	1.000000
<b>50%</b>	5.000000	1.993983	0.266375	3.315316	2.000000	1.000000
<b>75%</b>	6.000000	4.197569	0.388551	5.414627	3.000000	1.000000
<b>max</b>	14.000000	20.290516	0.844939	24.796182	9.000000	1.000000

```
# Trying to interpret 'Bounce Rate' by comparing it with Page Views
```

```
x = df[['Page Views']]
```

```
y = df[['Bounce Rate']]
```

```
plt.figure(figsize=(10,6))
```

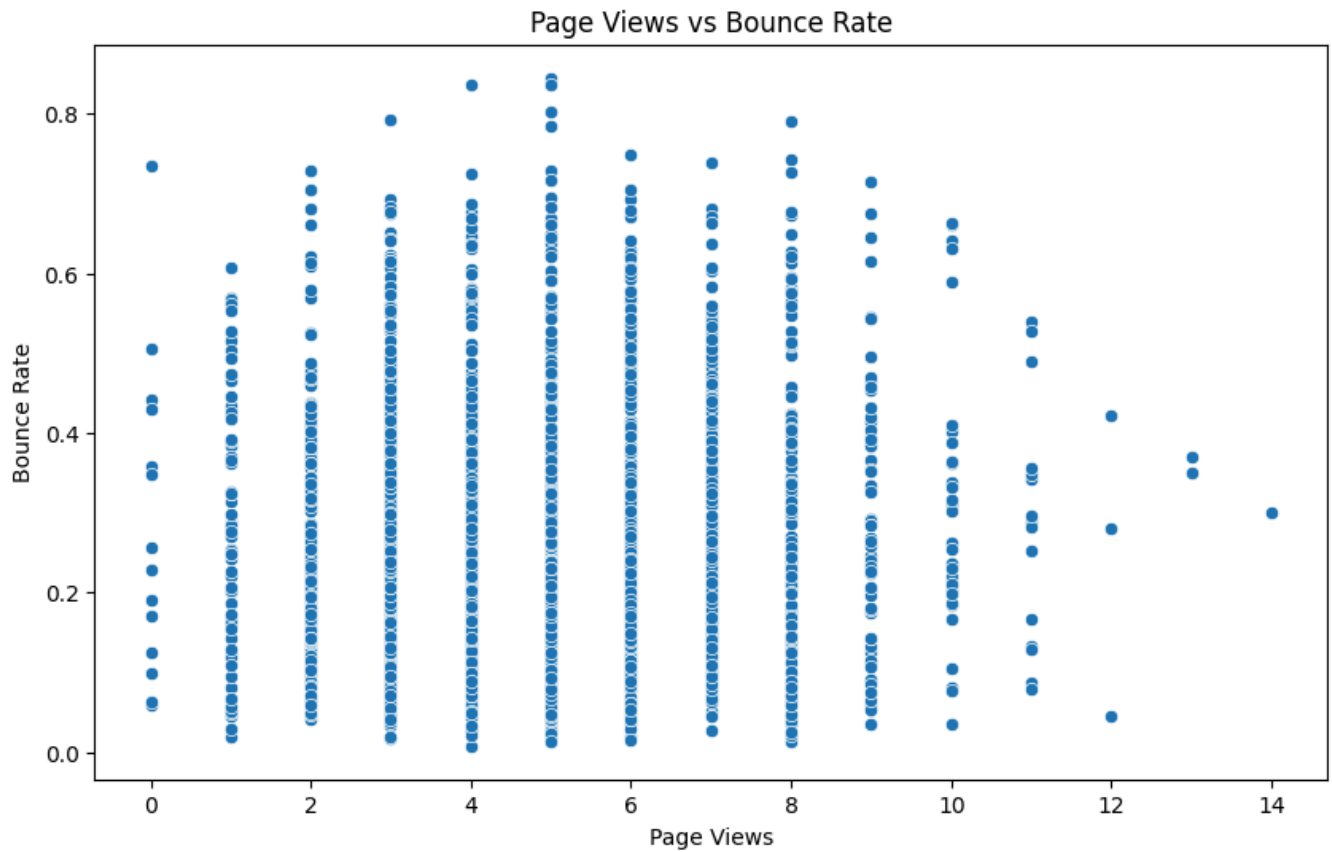
```
sns.scatterplot(x='Page Views', y='Bounce Rate', data=df)
```

```
plt.title('Page Views vs Bounce Rate')
```

```
plt.xlabel('Page Views')
```

```
plt.ylabel('Bounce Rate')
```


```
plt.show()
```




```
for col in df.columns:
    plt.figure(figsize=(10,6))
    sns.histplot(df[col], kde=True)
    plt.title(f'Distribution of {col}')
    plt.xlabel(col)
    plt.ylabel('Frequency')
    plt.show()
```



Distribution of Page Views

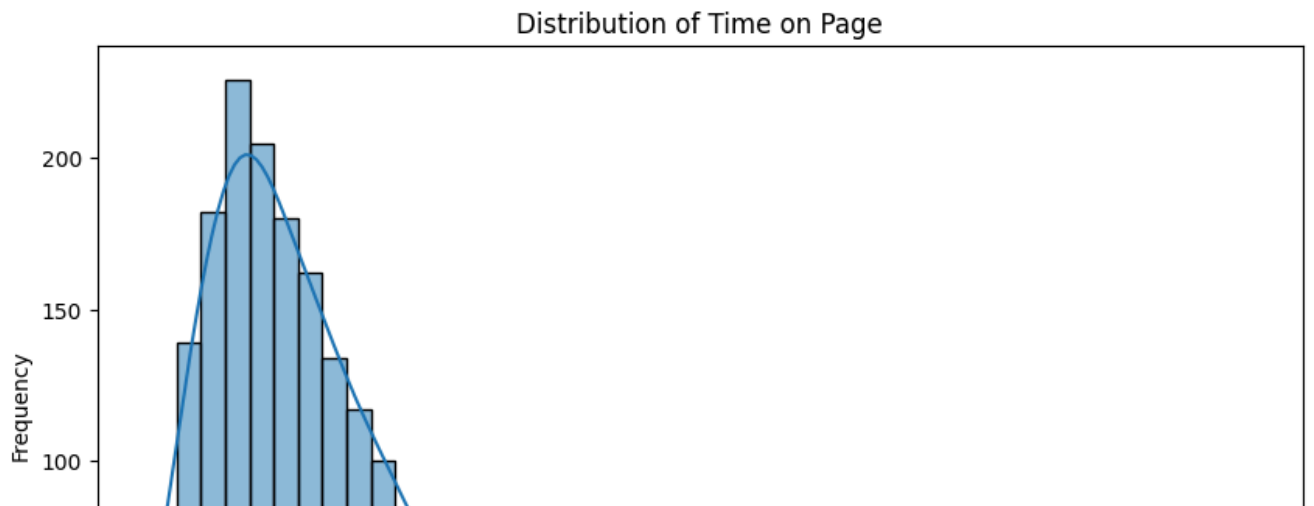





Distribution of Session Duration




Distribution of Bounce Rate









Distribution of Previous Visits



Distribution of Conversion Rate








It appears several features have a right-tail skew. The majority of data is clustered near the bottom of the data, but several possible outliers may be changing its shape. Features such as 'Session Duration', 'Previous Visits', 'Time on Page', and others all share this quirk.

```
df_numeric = df.select_dtypes(include=[np.number])

correlation = df_numeric.corr()


plt.figure(figsize=(10,6))
sns.heatmap(data=correlation, cmap='coolwarm', annot=True)
plt.title('Correlation Matrix (Numeric)')
plt.show()
```



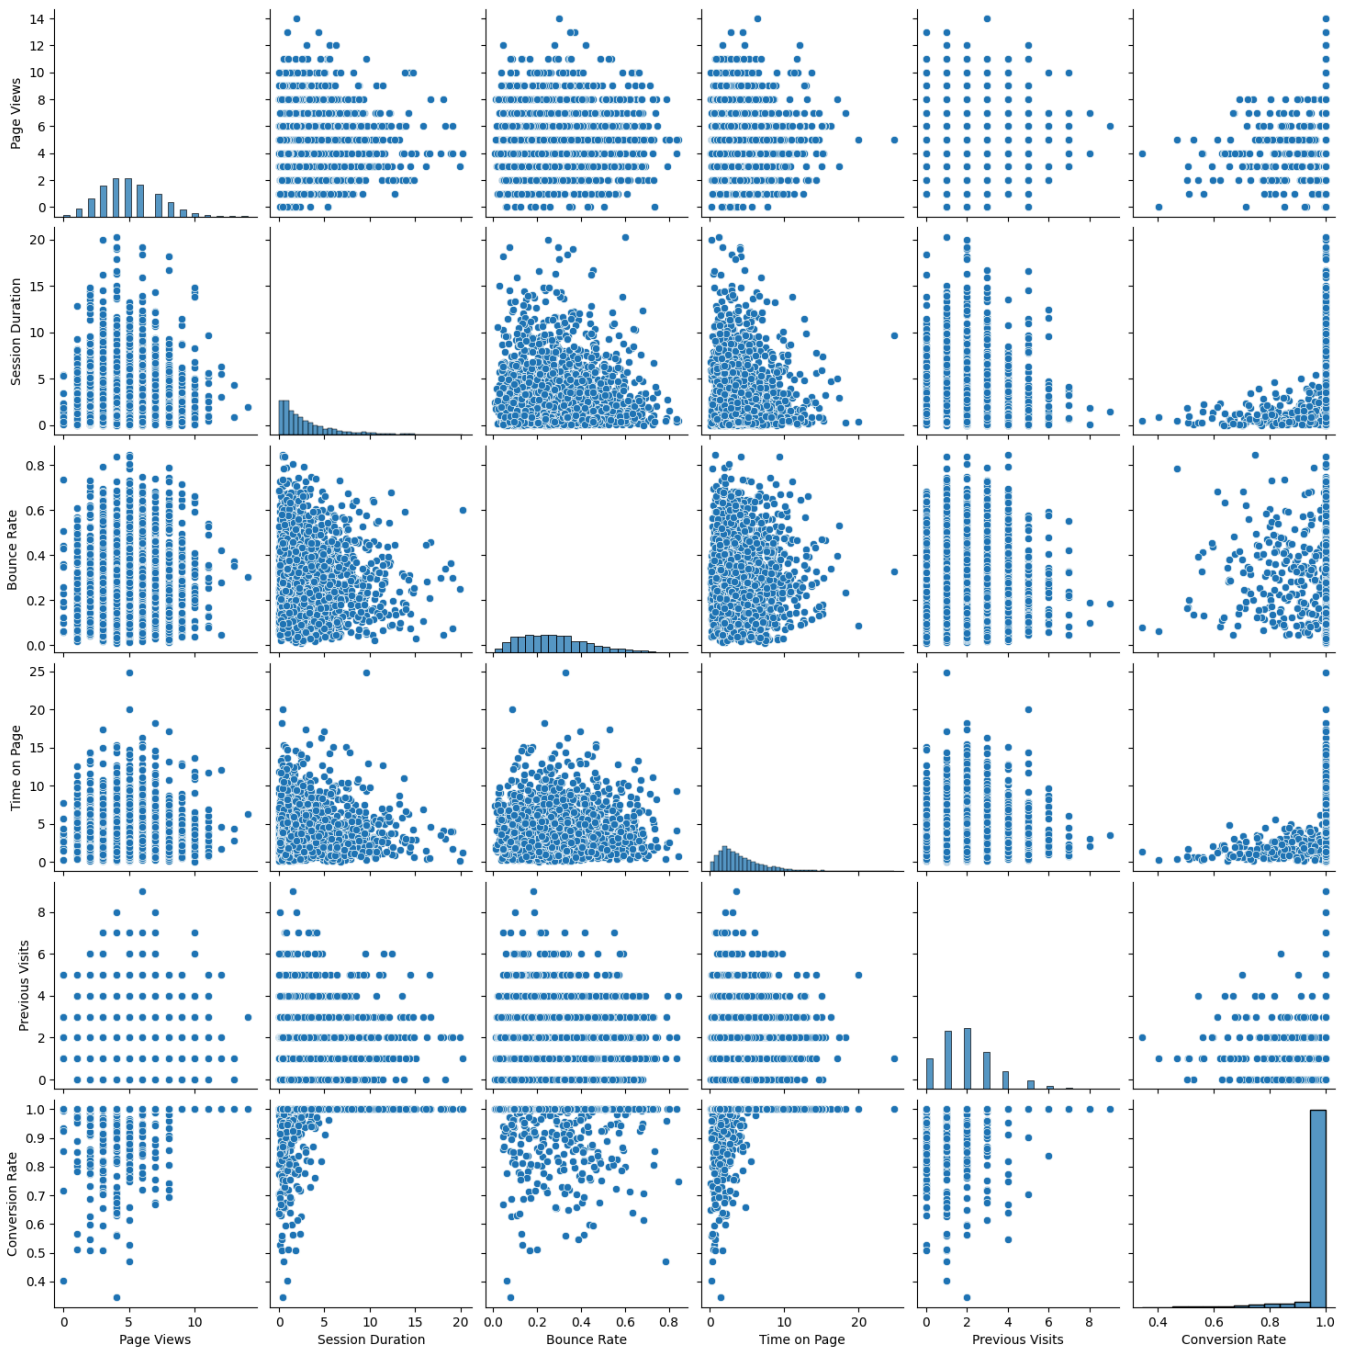
```
# One-hot encoding traffic type for correlation analysis
df_encoded = pd.get_dummies(df, columns=['Traffic Source'])
```

```
correlation_encoded = df_encoded.corr()
```

```
plt.figure(figsize=(10,6))
sns.heatmap(data=correlation_encoded, cmap='coolwarm', annot=True)
plt.title('Correlation Matrix (Encoded)')
plt.show()
```



```
# Pairplot for feature analysis
sns.pairplot(df)
plt.show()
```



All features in the dataset appear to have very low correlation with one another. Combined with all of the right-tail skews in the features, I suspect this may be due to the presence of outliers in the data. I will perform z-score and Inter Quartile Range analysis to test this.

```
# Finding outliers through z-score
df_outliers = df.copy()
outliers = []

for col in df_outliers.columns:
    if df_outliers[col].dtype in ['int64', 'float64']:
        mean = df_outliers[col].mean()
        std = df_outliers[col].std()
```

```
z_scores = (df_outliers[col] - mean) / std
```

```
# Printing all values with |z-score| > 3
for value in z_scores:
    if abs(value) > 3:
        outliers.append(value)
```

```
print(outliers)
```

```
[-3.5454072159846204, -5.431979194201322, -3.0674235923700857, -6.901097028876236, -7.21
```

```
numeric_columns = df.select_dtypes(include=[np.number])
outliers_iqr = find_outliers_iqr(numeric_columns)
print(outliers_iqr)
```

```
[11.05138124  9.63616963  9.68859192 13.58023187 10.02890318 18.3366796
 8.79349225 11.        11.69293807 11.65038153 15.0222197 10.
 9.00440239  9.        11.        15.0497133  9.23837253 14.33851631
 9.84027064 10.30849747 18.23996376  9.        9.        9.
 8.80059289 10.        9.65757692  9.        11.40041688  9.0817998
 9.2274092 11.62537824  9.64423172 10.67944941 10.97709127  9.
14.26480293  9.79144119 12.26653769  8.89755114  9.04986825  9.12855873
11.3484583 10.        9.49737922 10.        10.        13.33602859
11.        11.65479442 10.        9.02713664  9.16915655 16.28439777
10.9421283 13.54995784 20.0212847  9.72149194 10.75544236 10.
10.        13.63962071  9.        13.26421013 10.21244597  9.15937315
10.41547871 10.        14.33777637 12.2619853  9.        9.0695977
 9.05531257  9.        11.34137139 10.38416239 14.49940368 10.52734538
 9.59139401 12.00702784 10.        13.77601687 11.02958591 13.36763409
12.80147015  9.        9.55999952  9.        8.95568161 10.
11.82438278 10.52541519 10.45533011 10.        12.72859359  9.40939513
12.96431643 16.32290146  9.        12.        12.08349527 10.
 9.        9.54736229  9.        11.        9.13028533 19.93267164
10.62888345 11.62391698 10.28372383 10.73116579 11.43598842 11.
12.80008785  9.2959347  9.5483859  9.03304007 19.14363617  9.39737406
 9.        15.27774418  9.1816104 10.        13.8486278  9.
 9.        11.        9.        9.03133694 14.55003526 11.
12.13369056  9.70128077  9.40835794 13.68771409  8.78233322 12.88332128
11.        10.        11.        10.71673503 11.        10.
10.26290004  9.        9.17058449 10.        8.86475458  9.2460313
13.68943989 10.        9.9276564  9.        9.        13.61465526
10.78251722 10.        8.80277839 14.45754852  9.96589233  9.
 9.        9.        9.        10.43653859 10.        8.96751564
```

```

9.          9.          9.          9.          10.          11.07541224
10.         12.40890476  9.3310172  9.13607213  9.          13.88265279
9.          9.7659436  9.          10.56437726  9.16624747  10.21838447
11.         9.6449074  15.48708465  9.53525971  9.          8.88604007
9.21501953  11.51306737  9.          9.59640279  10.29315418  16.67657801
9.          11.52798876  10.48324902  10.6893614  9.          10.86486752
18.9563734  12.09657137  9.97781544  9.          10.98585239  8.84950476
10.70124688  11.81822948  10.11134025  9.58306823  10.39256813  11.34073594
9.19154493  9.          13.1017075  9.21772402  9.          9.20957212
11.94223739  20.29051597  10.94155166  11.77041037  10.          9.02040223
10.87521453  10.40772232  9.48149627  18.16045185  9.33093606  10.4732883
9.56802055  13.23233754  8.90209302  15.92677174  9.17579809  9.13277539
9.32370065  9.50834245  15.11562814  9.48255143  9.          10.70188101
9.29714524  9.57557761  16.55767846  14.31739623  12.03847918  10.12761527
9.          11.07889148  9.83254018  12.91064307  11.73366226  10.04452973
9.56051428  9.5669787  9.31248754  9.93669039  10.57907056  8.78230128
10.92763109  9.4318925  9.          15.10229463  10.6440976  12.
9.          9.72623619  9.          10.          9.28364265  9.47916263
14.68715848  17.83503341  10.          17.41481753  10.51224027  10.37711062
12.5211179  13.28557964  14.          8.90302423  9.63620996  9.27914382
12.64307642  9.65124136  24.7961822  11.          9.          10.58489758
13.99005495  11.13254018  10.          9.          14.30578208  9.
10.          11.05834852  12.72992129  10.77966016  9.74399319  15.0461243
9.          12.92470262  9.2171389  10.73565577  17.13063475  9.69535459
11.61178949  8.85497598  9.18913468  9.82909176  8.98928422  10.388859
11.50898232  12.14598649  9.62006223  12.07878624  10.          10.88130389
12.39653766  10.          14.63294927  19.1369552  9.          14.10801104
10.73190312  9.          13.          9.          10.75461815  11.
12.37674842  13.          8.80489367  10.34502843  11.33579979  9.
11.20004500  12.62000000  10.00000000  10.00000000  0.00000000  0.00000000

```

Based on EDA, it appears that there are no null values in the dataset. However, several features appear to be skewed, and there are multiple outliers in the data. These will need to be removed before analysis for accurate results.

## PREPARING THE DATA

```
from scipy import stats
```

```
# Removing outliers
```

```
# Computing Z-scores for all numeric columns
```

```
z_scores = np.abs(stats.zscore(df.select_dtypes(include=[np.number])))
```

```
# Keeping rows where all Z-scores are below the threshold (e.g., 3)
```

```
df_no_outliers_z = df[(z_scores < 3).all(axis=1)]
```

```
df_no_outliers_z.shape
```

```
➡ (1849, 7)
```

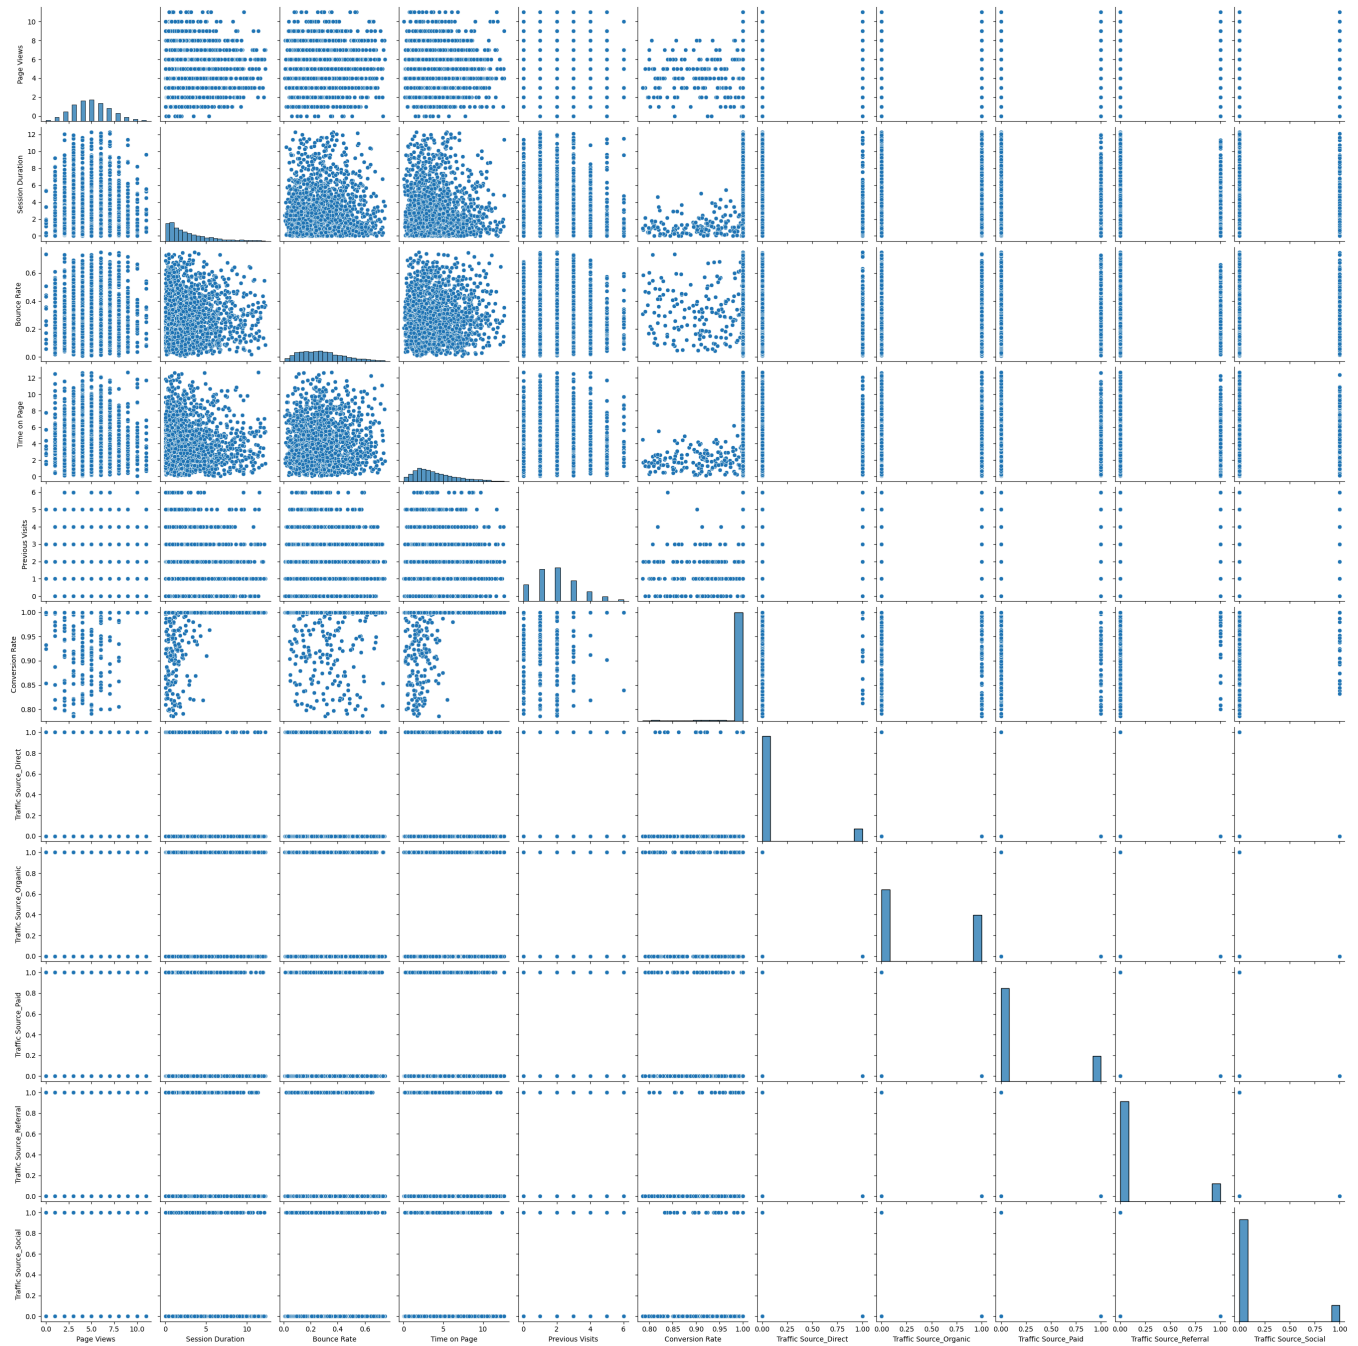
Removing outliers changed the number of rows from 2000 even to 1849, a drop of 151 rows.

```
# One-hot encoding 'Traffic Source' for linear regression
df_encoded = pd.get_dummies(df_no_outliers_z, columns=['Traffic Source'])
df_encoded.head()
```



	Page Views	Session Duration	Bounce Rate	Time on Page	Previous Visits	Conversion Rate	Traffic Source_Direct	Traffic Source_Organic
0	5	11.051381	0.230652	3.890460	3	1.0	False	True
1	4	3.429316	0.391001	8.478174	0	1.0	False	False
2	4	1.621052	0.397986	9.636170	2	1.0	False	True
3	5	3.629279	0.180458	2.071925	3	1.0	False	True
4	5	4.235843	0.291541	1.960654	5	1.0	False	False

```
# Pairplot for feature analysis w/o outliers
sns.pairplot(df_encoded)
plt.show()
```




The pairplot without outliers shows much more interesting information. Page views and bounce rate become much less skewed, and the connections between features and "conversion rate" become more pronounced - particularly "session duration", "bounce rate", and "time on page".

```

numeric_columns = df_encoded.select_dtypes(include=[np.number])
correlation_encoded = numeric_columns.corr()

plt.figure(figsize=(10,6))
sns.heatmap(data=correlation_encoded, cmap='coolwarm', annot=True)
plt.title('Correlation Matrix (Encoded)')
plt.show()

```



With outliers removed, some values gain more of a correlation:

- Session Duration has a (albeit very low) positive correlation with Conversion Rate.
- Time on Page also has a (albeit very low) positive correlation with Conversion Rate.

## MACHINE LEARNING